Module 3
Final Project

Will Dougherty
Flatiron School - Data Science - Online Self-Paced

Predicting Water Pump Functionality In Tanzania

Dataset:

- 74,250 water pumps
- 38 features

Tanzania:

- 60 million residents
- Over 1 million km^2

Problem:

- Limited resources in a large country
- Up-to-date functionality data is hard to come by

Project Goals:

- Predict **Non-Functionality**
- Use model to show how and where to efficiently allocate resources

Methodology

Focus on **Recall**

- Maximize % of non-functional pumps correctly identified
- Minimize false positives
  - mis-identifying *non-functional* pumps as *functional*

Build a **binary classification** model - ensemble of decision tree/random forest/etc

Main Data Problems

1: Missing-ness - ‘construction_year’

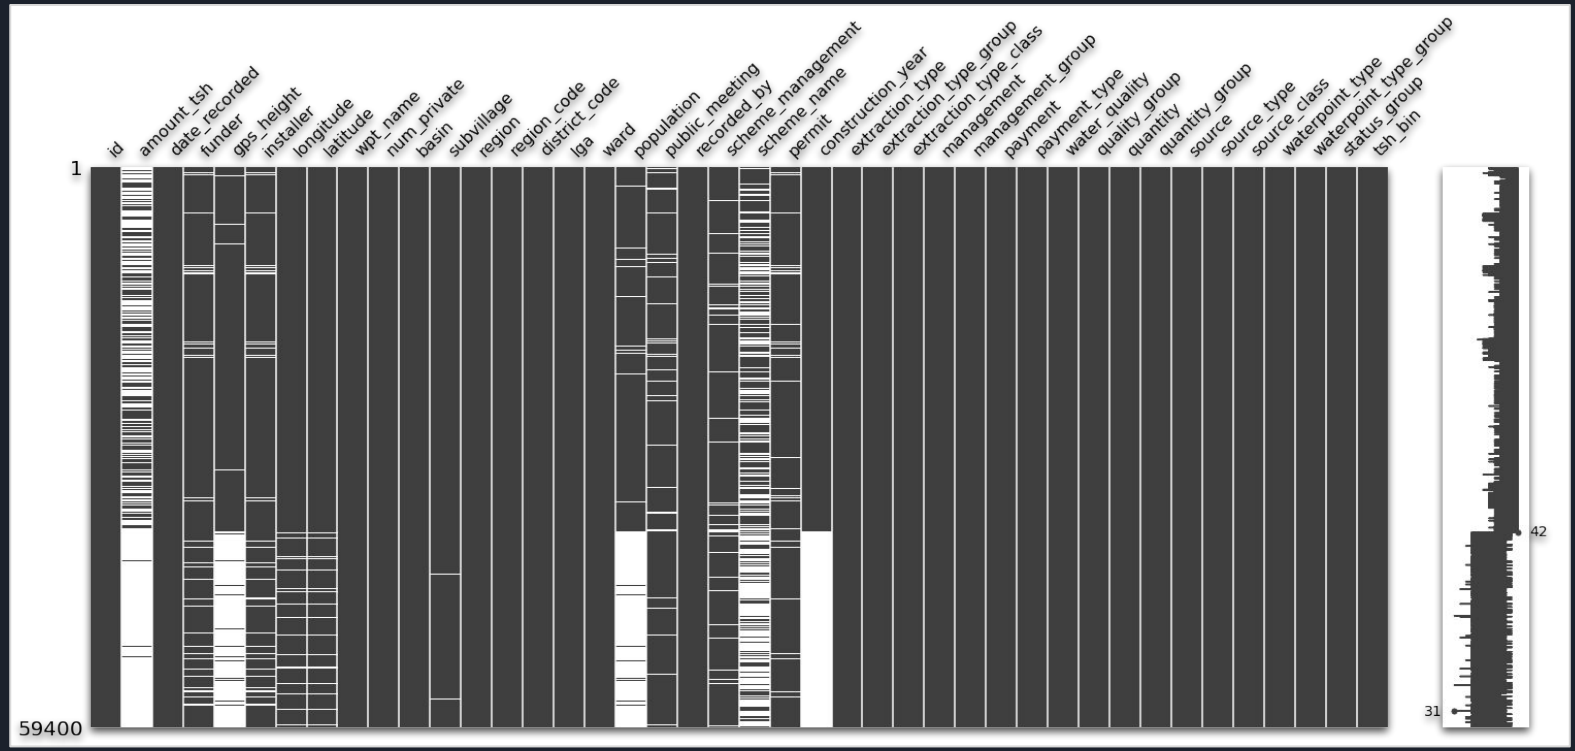

In pumps missing ‘construction_year’:

- Lack ‘amount_tsh’, ‘funder’, ‘installer’, ‘gps_height’, ‘longitude’/’latitude’, and ‘population’ in greater proportion

- Right: They are heavily clustered geographically

- Below: ‘age’ feature correlates strongly with functionality

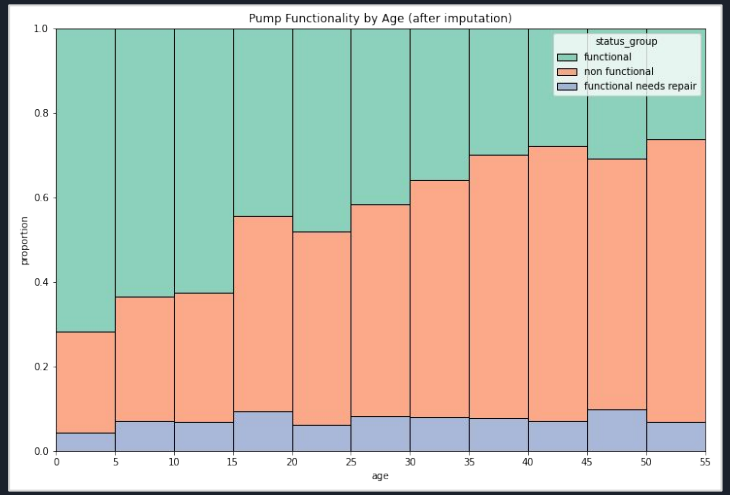

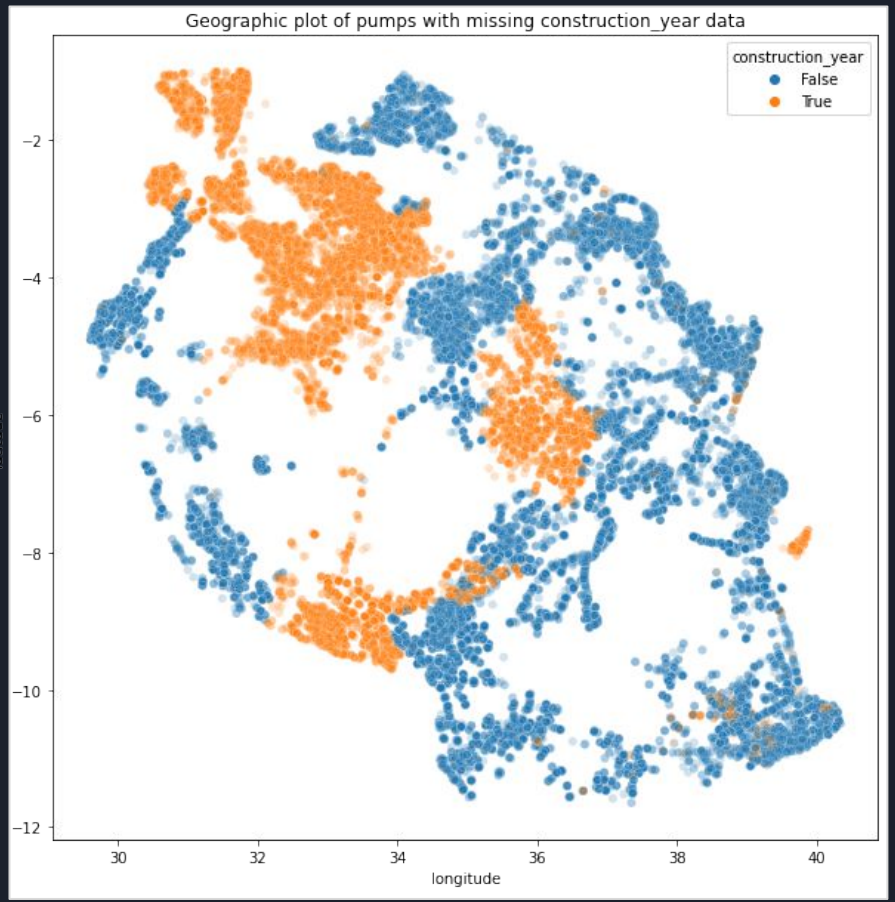

2: Non-standardization

‘funder’ and ‘installer’ features have thousands of unique values

Many appear to be misspelled or have inconsistent spaces/punctuation

Right: highest-similarity installer names

```
{'danida || danid': 0.909,
 'community || communit': 0.941,
 'gover || govern': 0.909,
 'tasaf || tassaf': 0.909,
 'fini water || fin water': 0.947,
 'oxfam || oxfarm': 0.909,
 'kiliwater || kili water': 0.947,
 'kiliwater || kilwater': 0.941,
 'rc church || rc churc': 0.941,
 'water aid || wateraid': 0.941,
 'consulting engineer || consuting engineer': 0.973,
 'muwsa || muwasa': 0.909,
 'finwater || fin water': 0.941,
 'villa || villag': 0.909,
 'fin water || finn water': 0.947,
 'adra/community || adra /community': 0.966,
 'adra/community || adra/ community': 0.966,
 'adra /community || adra/ community': 0.933,
 'local  technician || local technician': 0.97,
 'water aid /sema || water aid/sema': 0.966,
 'jandu plumber co || jandu plumber  co': 0.97,
 'muwasa || mtuwasa': 0.923,
 'tuwasa || mtuwasa': 0.923}
```

The Model

Ensemble classification model:

- Decision tree, Random forest, Bagging, XGBoost in a Voting Classifier

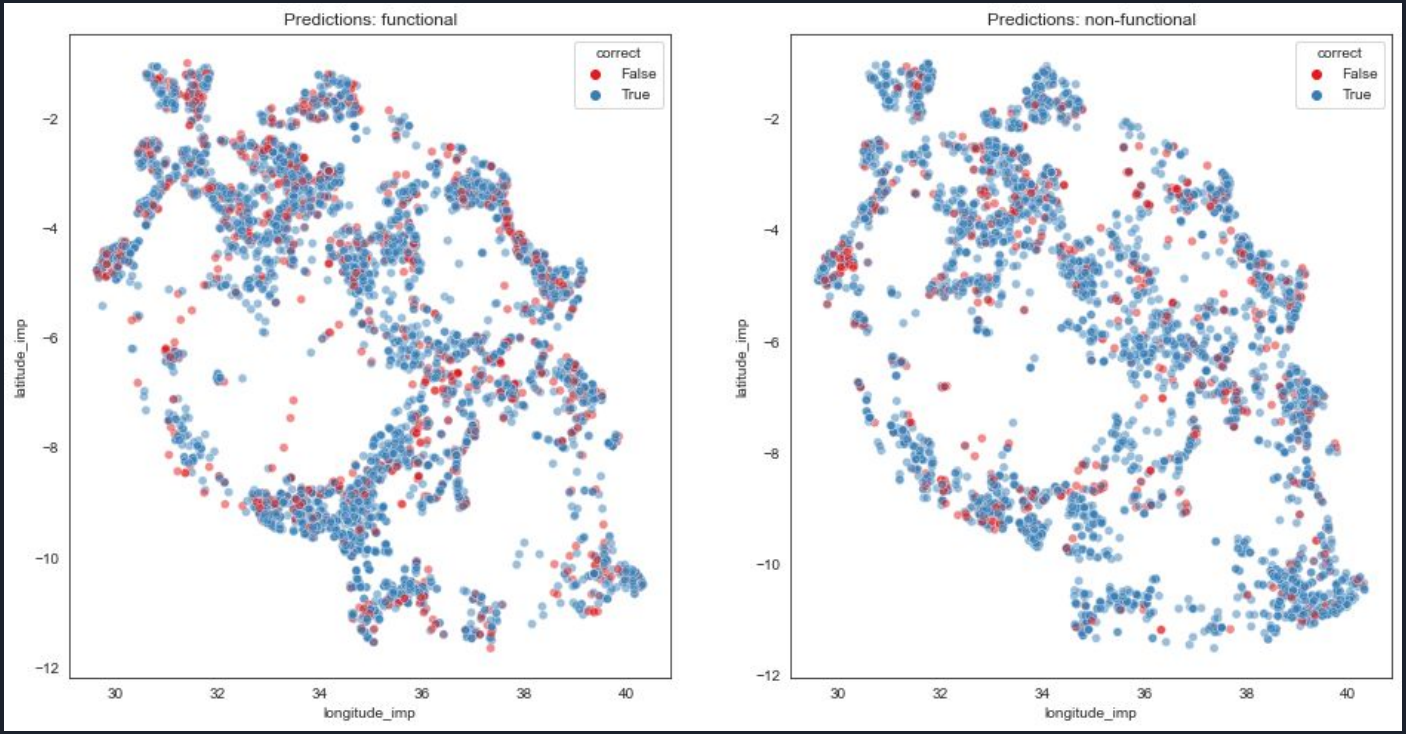

Recall: 85% of non-functional pumps are correctly identified

Regional Prediction

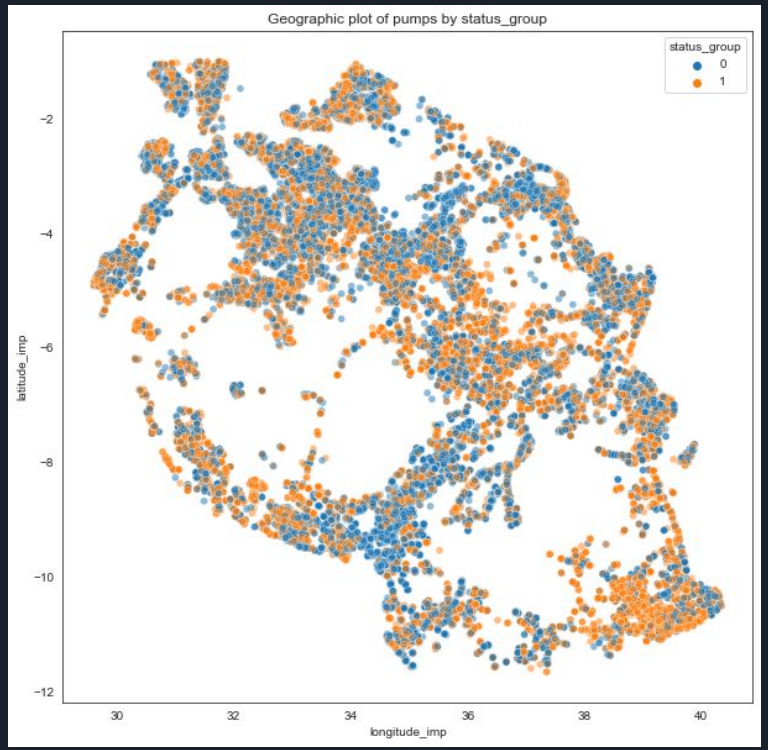

Functionality is clustered throughout the region

Knowing about individual pumps is less useful than knowing about areas/regions

‘Lga’ is a good middle-ground of granularity between large areas and villages

Tentative Results

The vast majority are predicted very accurately, closely packed around -0.1 - 0.0

The model tends to under-estimate rate of non-functionality compared to the actual rate

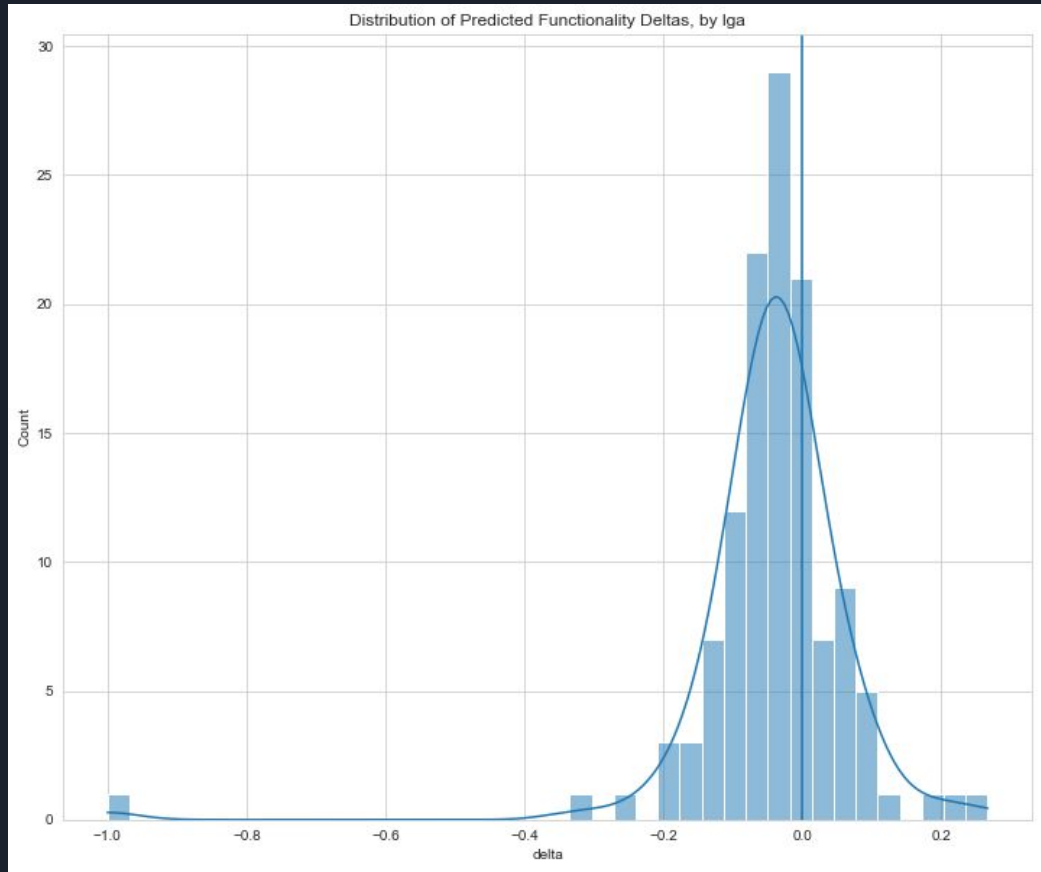

Map of non-functionality rates: Actual vs. Predicted

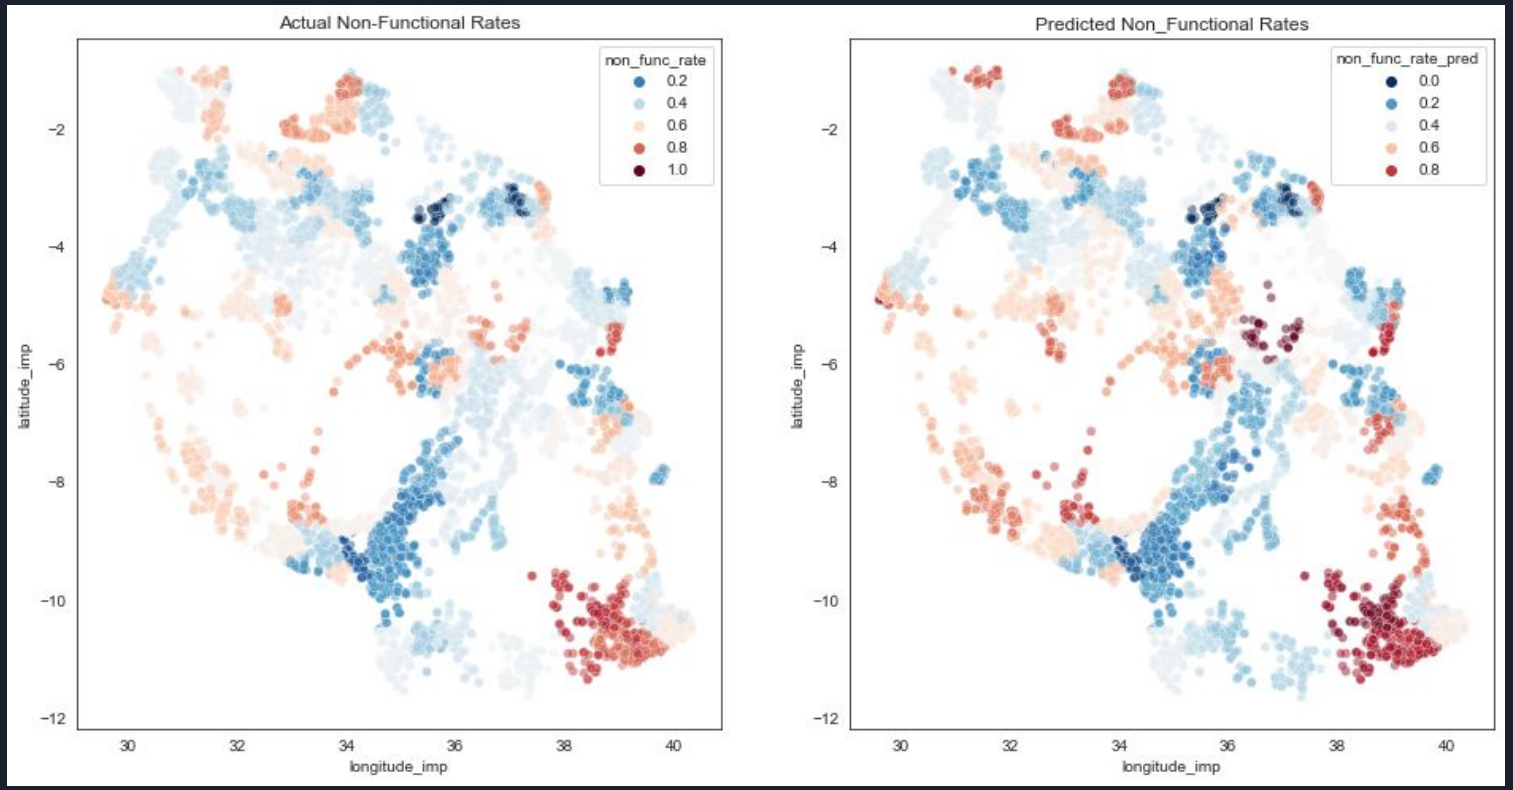

Conclusions

- The model does well in the **recall** metric, by correctly predicting 85% of the non-functional pumps

- Using the model, we should focus on **regional prediction**
  - Using ‘lga’, we can accurately predict rates for areas
  - This should be used to determine areas most in need of resources

Future Work

- Improve ‘lga’ and regional prediction
  - More consistent distribution of pumps - consolidate areas with few pumps
  - Determine population within each area to make resource allocation proportional

- Determine quality of model over time
  - Predict 5 years in the future, and test those pumps in 5 years to see how it performs

Thank you!